

**ROOM OCCUPANCY SENSING USING A
THERMAL TRIPWIRE**

Janis Intoy and Emily Lam

May 5, 2017

Boston University

Department of Electrical and Computer Engineering

Technical Report No. ECE-2017-01

**BOSTON
UNIVERSITY**

ROOM OCCUPANCY SENSING USING A THERMAL TRIPWIRE

Janis Intoy and Emily Lam



Boston University
Department of Electrical and Computer Engineering
8 Saint Mary's Street
Boston, MA 02215
www.bu.edu/ece

May 5, 2017

Technical Report No. ECE-2017-01

Summary

A large part of a building's efficiency depends on the efficiency of its HVAC (Heating Ventilation & Air Conditioning) system, which can be improved with automatic adjustments based on room occupancy level. As thus, accurate prediction of the number of people in a room is needed. In order to estimate a room's occupancy level, the Occusense Senior Design team has created a privacy-preserving, low resolution, thermal sensor system to capture temperature images of people walking through the doorways. This project uses that information coupled with a background prediction algorithm and optical flow analysis to predict the direction of motion of the people passing through and keep a continuous count of the number of people in the room.

We test our algorithm on people moving in and out, people lingering at the door, people rushing through the door, and multiple people passing together through the door. Our algorithm was not successful for people passing together through the door. Final results are 90% in classification and an error of ± 3 in the occupancy count after 10 people pass through the door in either direction.

Contents

1	Introduction	1
2	Literature Review	2
3	Problem Statement and Proposed Solution	2
3.1	Person Detection	2
3.2	Direction Discrimination	4
3.3	Classifier and Counter	4
4	Implementation	5
4.1	Data Acquisition	5
4.2	Person Detection	5
4.3	Direction Discrimination	6
4.4	People Count	7
5	Experimental Results	8
5.1	Background Subtraction	8
5.2	Direction Discrimination	9
5.3	People Count	12
6	Conclusions	12

List of Figures

1	Overarching system diagram. Our work focuses on the orange boxes.	1
2	Three-step implementation.	5
3	Example of detected foreground pixels.	6
4	Example of total number of foreground pixels in each frame compared to ground truth events in gray.	6
5	Example of optical flow vectors.	7
6	Mean vertical gradient for two people walking out.	7
7	ROC curves of detecting motion events using different parameters for the two proposed background subtraction algorithms.	8
8	Mean vertical gradient for lingering through the door.	9
9	Mean vertical gradient for running out and in through the door. . . .	10
10	Mean vertical gradient for two people walking in together	10
11	Proportion of correct direction classifications with peak performance at 90%.	11
12	Histogram of people count errors after 1 motion event (blue) to 10 motion events (yellow).	11

List of Tables

1	Good parameter ranges for background subtraction algorithms.	6
---	--	---

1 Introduction

In modern efficient buildings, controlling the HVAC system preemptively can be a huge cost saver. Therefore, Occusense, a senior design team at Boston University, is working towards developing sensing technology and algorithms to continuously keep count of the number of people in a room. This way, the system can adjust system parameters, such as ventilation, accordingly, before feedback sensing technology can detect abnormalities in the air quality of a room. There are a number of ways to detect occupancy, such as using a fisheye camera to count the number of people in a frame. However, this project assumes a tripwire methodology, that is if a room has a low number of entry points, counting people can be done at the entries based on if someone or something is entering or leaving the room. As such, continuous knowledge of the number of people in a room is achieved as a running tally. The senior design team has implemented a low resolution thermal sensor with a field of view of $30^\circ \times 120^\circ$ positioned at the top of the door frame and looking perpendicularly down, Figure 1. It is capable of capturing a 4×16 pixel array of temperatures at frame rates of 8-12Hz. Using this information, our project develops an algorithm to count the number of people entering or leaving a room. Specifically, this algorithm will 1) detect the presence of a moving person in the frame, 2) determine the direction of motion of the person, and 3) keep count of the total number of people in a room.




Figure 1: Overarching system diagram. Our work focuses on the orange boxes.

2 Literature Review

Many background subtraction algorithms have been developed to detect changing pixels in a series of images. A brief review of simple background subtraction methods can be found in [1] and more common change detection algorithms in [2]. The most successful amongst these incorporate a background model into the algorithm. The use of a model allows for thresholding of probabilities instead of pixel intensities and is therefore more robust to some variation in the background scene. In addition, foreground models can improve the sensitivity of the change detectors [3]. McHugh et al. suggested a foreground model algorithm that is more general as it is based on spatial neighborhoods. To further improve the discrimination, they also suggest a Markov model so that labels are more spatially coherent [4].

We also looked into algorithms for determining optical flow. Optical flow describes the direction elements in an image move or change with space and time from frame to frame, usually in a vector form. There are three overarching approaches for determining optical flow: a feature-based approach, a correlation-based approach, and a gradient-based approach. For our project we consider the gradient-based approach on foreground elements so that only the foreground elements are put through the gradient-based optical flow algorithm when determining object motion. While this method relies on background subtraction for foreground selection, it does not require prior training. The gradient-based optical flow algorithm models gradients based on partial derivatives for each pixel intensity with respect to space and time, as each pixel's intensity will have a change in both the x-direction, y-direction, and time [5].

3 Problem Statement and Proposed Solution

Our project aims to keep track of the number of people in a room based on the tripwire methodology. Our approach breaks the problem into three parts: 1) detect a person in the doorway, 2) determine the direction of motion of the person, 3) classify the motion and track the total number of people in a room.

3.1 Person Detection

In order to accomplish this we make several assumptions about the acquired data. First, we assume that the background has slow temporal dynamics. Second, we assume that the data starts with some number of background-only frames. For business and residential buildings this is likely true if they close for the night. Finally, we assume that the noise in the surface temperature measurements is independent of the signal. Here we propose two background subtraction methods to detect a person or persons in the frame. One method to detect foreground in each frame is to model the background temperatures of each pixel as a Gaussian which is updated with each new frame. The parameters of the background model are updated when the current

pixel is labeled as background:

$$\mu_t[\mathbf{n}] = \rho I_t[\mathbf{n}] + (1 - \rho)\mu_{t-1}[\mathbf{n}]$$

$$\sigma_t^2[\mathbf{n}] = (I_t[\mathbf{n}] - \mu_t[\mathbf{n}])^2 \rho + (1 - \rho)\sigma_{t-1}^2[\mathbf{n}]$$

where $I_t[\mathbf{n}]$ is the pixel temperature at position $[\mathbf{n}]$ and μ_t and σ_t^2 are the mean and variance of the gaussian in frame t . A Markov model can be applied to improve the spatial coherency of foreground labels: a pixel surrounded by foreground pixels has a low threshold since it is likely also foreground. Therefore, a pixel is labelled as foreground when

$$\frac{I_t[\mathbf{n}] - \mu_{t-1}[\mathbf{n}]}{\sigma_{t-1}[\mathbf{n}]} >_{\mathcal{F}} \theta \exp((Q_{\mathcal{B}}[\mathbf{n}] - Q_{\mathcal{F}}[\mathbf{n}])/\gamma)$$

where θ is an initial threshold ($\#$ of standard deviations), $Q_{\mathcal{B}}[\mathbf{n}]$ and $Q_{\mathcal{F}}[\mathbf{n}]$ are the number of background and foreground pixels neighboring the pixel at position \mathbf{n} respectively, and γ is a parameter that changes the influence of the Markov model on the threshold.

The second proposed solution is inspired primarily by the foreground-adaptive background subtraction described in [4]. The background can be modeled by a probability density function estimated at each pixel location \mathbf{n} from the N most recent pixels that were labeled as background:

$$P_{\mathcal{B}}(I_t[\mathbf{n}]) = \frac{1}{N} \sum_{i \in \mathcal{B}_t[\mathbf{n}]} \mathcal{K}(I_t[\mathbf{n}] - I_i[\mathbf{n}])$$

where $\mathcal{B}_t[\mathbf{n}]$ are the N previous time indices at which the pixel located at \mathbf{n} was labeled background, and \mathcal{K} is a zero-mean Gaussian with variance σ^2 . Thresholding these probabilities with threshold θ gives an initial classification of pixels into background or foreground.

The development of the foreground model follows a similar formulation, but the summation is over pixels in the neighborhood of \mathbf{n} that already have been labeled as foreground. The ratio of the probability density functions $P_{\mathcal{B}}$ and $P_{\mathcal{F}}$ can be thresholded to determine new labels for each pixel. Furthermore, the number of foreground neighbors a pixel has can be used to adjust the threshold. The more foreground neighbors a pixel has the easier it is to be labeled as foreground. This will contribute to the spatial coherence of the labels. Thus, a pixel is labeled as background if

$$\frac{P_{\mathcal{B}}(I_t[\mathbf{n}])}{P_{\mathcal{F}}(I_t[\mathbf{n}])} > \theta \exp\left(\frac{1}{\gamma}(Q_{\mathcal{F}}[\mathbf{n}] - Q_{\mathcal{B}}[\mathbf{n}])\right)$$

where γ is a parameter to control how strongly the threshold adapts. The thresholding step can be done iteratively as labels are updated in each step.

3.2 Direction Discrimination

After determining the foreground from the background, we look at determining the direction of flow of the foreground elements. We do this by taking the foreground elements only, setting the background elements to zero, and applying a gradient-based optical flow algorithm. This can be modeled as described in [5]:

$$\frac{\partial I}{\partial x} \frac{\partial x}{\partial t} + \frac{\partial I}{\partial y} \frac{\partial y}{\partial t} + \frac{\partial I}{\partial t} = 0$$

Since we are only concerned with the vertical velocity, i.e. a person walks through the door frame and the laterally motion is trivial, the previous equation can be written as

$$v_y = -\frac{\frac{\partial I}{\partial t}}{\frac{\partial I}{\partial y}}$$

where $v_y = \frac{\partial y}{\partial t}$. A large v_y would correspond to a lot of motion in that direction. In our system of 4x16 array of pixels, we generate a spatial gradient matrix consisting of 3x16 pixels to take the boundary condition into account. Next, we perform two averages, we average the rows to find the overall spatial gradient of the row, and we then average those row averages together to obtain an overall spatial gradient for the entire array per a time. Next, the temporal gradients are obtained between consecutive frames. The final velocity is then calculated and used to determine the optical flow of the foreground which determines whether a person is walking into or out of the room.

3.3 Classifier and Counter

A classifier determines when an event starts and stops based on the background subtraction algorithm. It then compares the spatial gradients at different times during each event and increments the overall counter accordingly. In addition, it takes into account edge cases for determining between noise and data for single-frame events.

4 Implementation

The current implementation of the algorithm is in Matlab. The data, code, and documentation is available on github: <https://github.com/janisirene/Occusense>. The three parts of the algorithm, Figure 2, are implemented separately and can be run in sequence to output a single room occupancy number. We define an event as a series of consecutive frames with detected motion as determined by the background subtraction algorithm.

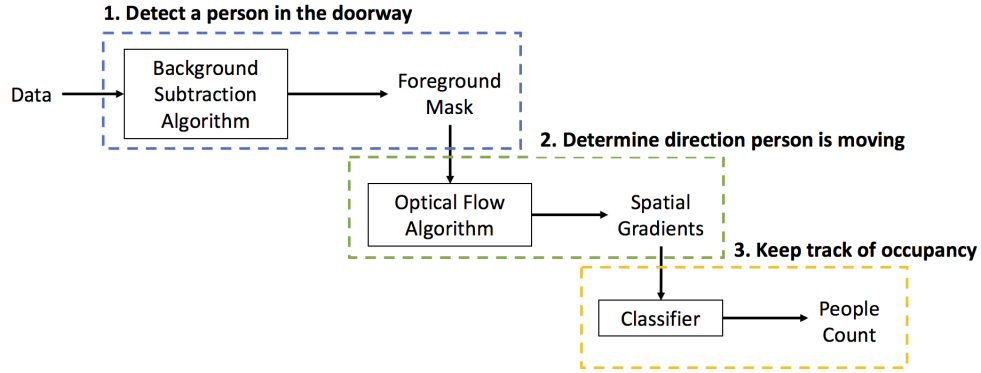


Figure 2: Three-step implementation.

4.1 Data Acquisition

In addition to the labelled data from the senior design team, we acquired several trials of realistic but more difficult situations. This includes a person lingering in the doorway, a person rushing out the door and back in, and multiple people passing through the door in quick succession.

The data was recorded into text files at a rate of 8-12Hz. The labeled data was recorded at 8Hz by the Occusense team and the difficult situations data was recorded at 12Hz.

4.2 Person Detection

Janis implemented both proposed background subtraction algorithms in background-SubtractionSimple.m (running gaussian) and backgroundSubtraction.m (kernel-based approach). They take as input a structure of parameters and a 3-D array of temperatures over the two spatial and one temporal dimensions. The best parameter ranges are described in Table 1. This implementation could be modified to run in real time and, depending on the size of the history for the computation of the background PDF, would not require too much memory. The output of this function is a binary array of

pixel labels (red dots in Figure 3). Motion events are identified by thresholding the total number of detected pixels (Figure 4).

Foreground Adaptive			Running Gaussian Average		
	Description	Good Range		Description	Good Range
N	# background frames	10-30	N	# background frames	10-30
σ	std of gaussian kernel	0.4-0.6	ρ	size of temporal window	0.01
γ	influence of MRF	0.2 - 0.8	γ	influence of MRF	> 1
N_η	neighborhood order	1	N_η	neighborhood order	1
it	labelling iterations	3	it	labelling iterations	3
			θ	threshold	2.5

Table 1: Good parameter ranges for background subtraction algorithms.




Figure 3: Example of detected foreground pixels.




Figure 4: Example of total number of foreground pixels in each frame compared to ground truth events in gray.

4.3 Direction Discrimination

Emily implemented this algorithm in the file `opticalflow.m`. It takes as input a 3-D array of temperatures over the two spatial and one time dimensions. This can be the raw information provided by the sensors or a masked foreground-only version that has

already undergone background subtraction. This function returns spatial gradients in various forms: for individual pixels (Figure 5), averaged across rows, and averaged per frame (Figure 6). This information is passed to pCounter.m to perform final temporal gradients, direction decisions, and track occupancy. This implementation could be modified to run in real time to follow the background subtraction function.




Figure 5: Example of optical flow vectors.




Figure 6: Mean vertical gradient for two people walking out.

4.4 People Count

Janis and Emily implemented this algorithm in the file pCounter.m. This file takes the output of the opticalflow.m, a series of vertical gradients over time, and averages them per a frame and calculates the temporal gradients based on the first and last time frame as the optical flow information for the frames corresponding to the middle of an event are too noisy. The direction is then decided based on this information. This function also decides when single-frame events are considered noise are not based on the mean number of foreground pixels. If there are less than 4 foreground pixels,

the frame is considered to be noise and discarded. In addition, this information is used to determine when an event starts and stops.

5 Experimental Results

Janis To Do: Better explain test data.

Algorithms were tested on simulated data generated by randomly concatenating the labelled data. Since all of the data came from the same setup where the sensor was on the inside of the room, data were also randomly vertically flipped to remove the directional bias. The temperatures of the frames are shifted for continuity between frames.

5.1 Background Subtraction

The background subtraction results are shown as ROC curves (Figure 7). The kernel-based method had overall better performance than the parametric method achieving hit rates of over 90% for the best parameter sets with negligible false alarms in detecting a person in the frames. Most misses occurred because of spontaneous rapid changes in temperature of background pixels. These occur infrequently in the actual data, but also arise when true foreground pixels are mislabelled and bias the background model, for example at the edges of the foreground person. The effects of this are seen in Figure 4, where the second peak in an event is sometimes much lower than the preceding one.

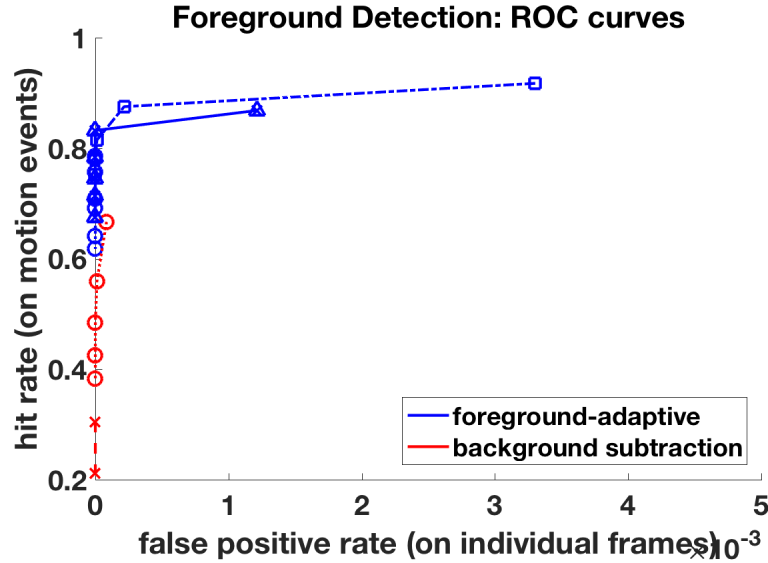


Figure 7: ROC curves of detecting motion events using different parameters for the two proposed background subtraction algorithms.

5.2 Direction Discrimination

The direction discrimination works for all cases we tried except for multiple people walking in together. For the simple case of two people walking out, Figure 6, the background subtraction algorithm correctly identifies two events, and the direction is easy to discriminate from the optical flow. For lingering at the door, Figure 8, because the noisy information in the middle frames is ignored, this also works well with our algorithm. In the case of someone running out and in quickly, Figure 9, again our algorithm is successful in determining the direction and successfully determines single frame events as motion versus noise. However, in the case where the background subtraction algorithm fails to detect individual people as individual events, the direction discrimination fails because it does not consider the middle frames in its decision, Figure 10. This could be remedied by taking into account the middle frames and employing peak counting or zero-crossing analysis. However, since the middle frames provide noisy information as seen in the lingering example, this information cannot be relied on when counting the number of people entering or leaving. Overall, the direction discrimination works well, Figure 11, and the major result in error is due to the error in the background subtraction algorithm.




Figure 8: Mean vertical gradient for lingering through the door.




Figure 9: Mean vertical gradient for running out and in through the door.




Figure 10: Mean vertical gradient for two people walking in together




Figure 11: Proportion of correct direction classifications with peak performance at 90%.




Figure 12: Histogram of people count errors after 1 motion event (blue) to 10 motion events (yellow).

5.3 People Count

Room occupancy is computed as cumulative sum of people entering and exiting the room. Since the errors here integrate, we find that the variance of the error is quite large after only several motion events. The standard deviation of the error after 10 events is up to 3 (Figure 12).

6 Conclusions

Although these algorithms perform well in capturing individual motion events, the resulting integrated people count is not a reliable estimate of room occupancy. There are several possible ways that this method could be improved. The first is to increase the frame rate of the thermal sensors to improve capture of fast motion and gaps between two people passing through the door in quick succession. This could also result in a different distribution of background noise but this can be dealt with by preprocessing the data or by the background model. Secondly, the detection of motion events could be improved to split single events that are suggestive of multiple people. For example, when two people pass through quickly, one event is labeled but there are two clear peaks in the number of detected pixels. Finally, this split could also occur in the direction detection state where zero crossings of the vertical gradient indicate a person passing through the doorway. In addition, logically the next step would be to implement the algorithms in real-time.

References

- [1] M. Piccardi. “Background subtraction techniques: a review,” *IEEE International Conference on Systems, Man and Cybernetics*, 2004.
- [2] N. Goyette, P. Jodoin, F. Porikli, J. Konrad, and P. Ishwar, “A Novel Video Dataset for Change Detection Benchmarking,” *IEEE Transactions on Image Processing*, vol. 23, pp. 4663-79, Nov. 2014.
- [3] A. Elgammal, R. Duraiswami, D. Harwood, and L. Davis, “Background and Foreground Modeling Using Nonparametric Kernel Density Estimation for Visual Surveillance,” *Proceeding of the IEEE*, vol. 90, pp. 390-393, July 2002.
- [4] J. McHugh, J. Konrad, V. Saligrama, and P. Jodoin, “Foreground-Adaptive Background Subtraction,” *IEEE Signal Processing Letters*, vol. 16, pp. 390-393, May 2009.
- [5] S. Smith, “Reviews of Optic Flow, Motion Segmentation, Edge finding and Corner Finding,” *Technical Report TR97SMS1, Oxford Centre for Functional Magnetic Resonance Imaging of the Brain (FMRIB)*, <https://users.fmrib.ox.ac.uk/~steve/review/review/review.html>, 1997.